NTRL
Information Architecture
Canonical Volume

Authoritative reference for categorization, feed structure, admin control, and system governance.
This document supersedes all prior IA-related specs.

Version 1.0 · Production System

PART I — Core Information Architecture

This section defines how NTRL understands, categorizes, and structures information.

These rules apply globally and are not UI-dependent.

1. Feed Categories (User-Facing)

Feed categories are stable nouns that users explicitly select.

Available categories include: World, United States, Local, Business, Technology, Science, Health, Environment, Sports, Culture.

Categories without qualifying stories are silently omitted.

2. Internal Domain Taxonomy (System-Only)

Each article is assigned exactly one internal domain. Domains are exhaustive and non-overlapping.

- Global Affairs
- Governance & Politics
- Law & Justice
- Security & Defense
- Crime & Public Safety
- Economy & Macroeconomics
- Finance & Markets
- Business & Industry
- Labor & Demographics
- Infrastructure & Systems
- Energy
- Environment & Climate
- Science & Research
- Health & Medicine
- Technology
- Media & Information Systems
- Sports & Competition
- Society & Culture

- Lifestyle & Personal Affairs
- Incidents & Disasters

3. Domain → Feed Category Mapping

Internal domains map deterministically to feed categories via admin-controlled rules.

This indirection preserves accuracy while keeping the UI calm.

PART II — Categorization & Feed Construction

This section defines how articles are classified and assembled into the feed.

4. Canonical Tagging Axes

All articles are tagged across fixed, non-evaluative axes.

Axes

Domain · Geography · Actors · Action Type · Temporal State · Evidence State · Impact Scope · Quantitative Signals.

Tags may not encode urgency, importance, moral framing, or political leaning.

5. Categorization Pipeline

Decision order: Geographic rules → Domain classification → Domain mapping → Conflict resolution → Confidence scoring → Review flagging.

All categorization outputs are explainable.

6. Feed Assembly Rules

Feed assembly respects user-selected categories and ordering.

Articles are sorted by time within category.

Empty categories are omitted entirely.

PART III — Admin Control & UX

This section defines how humans safely control and correct the system.

7. Admin Capabilities

Admins may manage feed categories, domain mappings, quotas, and overrides.

Low-confidence categorizations flow into a structured review queue.

8. Override & Audit Rules

Overrides require justification and are immutable.

All changes are logged and versioned.

PART IV — Machine Contracts & Governance

This section defines the machine-readable backbone of the system.

9. Taxonomy Versioning

All taxonomy artifacts are versioned.

Articles store the taxonomy_version used at classification time.

```
{
  "taxonomy_version": "1.0.0",
  "effective_date": "2026-01-27"
}
```

10. Admin API Contract

The Admin API is the sole interface for mutating categorization behavior.

No silent changes are permitted.

```
PUT /admin/domain-mapping
POST /admin/override
GET /admin/review-queue
```

11. Rollback Semantics

Taxonomy versions may be rolled back forward-only.

Historical content is not silently reclassified.

12. Design Outcome

This architecture enables scale, auditability, and long-term neutrality.

It prevents entropy, editorial drift, and manipulation by design.